

Project practice

2024/2025

Extending Whisper for Target Speaker ASR

Daniel Bohata*

Abstract

The problem of transcribing the speech of a single person in multi-speaker environments is a complex one with some existing solutions. In this work I have sought to replicate and attempt to improve one such approach. The chosen approach is extending a well known single speaker transcription model (Whisper) for this task using prompt tuning [3]. That means working with a frozen model by injecting soft prompts into the encoders, possibly decoders as well. My results have been achieved by replicating a paper by the authors of this method [3]. Furthermore, I analyzed the results when applying more training epochs, using different schedulers and mixing both clean and noisy datasets together. Training and evaluation has been done using GPU clusters provided by MetaCentrum. I have obtained results with an average difference of 1.12% from the original. Further experiments have not shown significant improvement over previous results. My results are within minimal margin from results claimed by the authors of this approach and my experiments did not succeed in improving them.

Keywords: Target Speaker — Automatic Speech Recognition — OpenAI Whisper

Supplementary Material: [Original code](#) — [Improved code](#)

*xbohatd00@fit.vut.cz, Faculty of Information Technology, Brno University of Technology

1. Introduction

Target speaker automatic speech recognition (TS ASR) is a well-known problem with many use cases. First, it has been solved by training a model from scratch. Then, fine-tuning single speaker models came about, but that takes a lot of time and resources. In this work, I aim to reproduce results of another method. The method of choice is extending an already trained single speaker model, OpenAI's Whisper, for this task using prompt tuning. This should save both time and resources required to achieve transcription. The method of choice has already been studied and published with results; hence, I am only aiming to reproduce the results claimed and attempt to improve them.

The aim of TS ASR is to accurately transcribe speech from a specific target speaker in an audio that may contain overlapping speakers or background noise. The core problem is that existing state-of-the-art methods either train a model from scratch or fine-tune an already existing one, both of which are time-consuming

and resource-consuming [2]. The correct solution should ideally decrease both the time and the resources needed to successfully complete the transcription as quickly as possible while minimizing the word error rate (WER).

A solution to this already exists and is the core of this work. It is based on expanding Whisper's capabilities using prompt tuning instead of fine tuning. The code and the paper have been published by Hao Ma, Zhiyuan Peng, et al. [2]. Their code is used as the basis for my implementation [3].

My solution is to expand on the existing implementation by achieving better results eg. lower word error rate. The aim is to do that by using different schedulers, longer training, or training on a combined dataset. Unfortunately, the code and instructions that came with it had some problems that had to be solved. First, the parameter *experiment name* was actually the name of checkpoints that would be created in training. Second, there was a type conversion error in the

user	fairshare	job count					
		m1	total	queued	running	completed	other
tasevsky	189	10001	0	0	0	10001	
galaxyelixir	188	102	0	8	94	0	
galaxyumsa	187	46	0	0	46	0	
mark_cheeky	186	3	0	1	2	0	
xveselsky	185	32	0	0	32	0	
lunak	184	1	0	0	0	1	
janavisoanova	183	1	0	0	1	0	
jendrb00	182	4	4	0	0	0	

Figure 1. Efficiency when running jobs is key, as we can see that someone with 10000 jobs has the highest fairshare

41 code that needed to be fixed or else training was not
42 possible.

43 2. Methodology

44 In this section, experimental setup, Whisper and prompt
45 tuning is introduced. Furthermore, the process of ob-
46 taining the training data set and speaker embeddings
47 is also introduced.

48 2.1 Experimental setup

49 To replicate the results, a powerful GPU was needed
50 as Hao Ma's group used a 24GB Nvidia 3090, which
51 I do not have. For that, the MetaCentrum cluster net-
52 work was chosen. A logging tool was also missing
53 in the code, for that weights and biases was chosen
54 and implemented. In this section, MetaCentrum, it's
55 scheduling system PBS and metrics logging will be
56 described.

57 2.1.1 MetaCentrum

58 MetaCentrum is a free platform for students and em-
59 ployees of Czech universities that allows students to
60 try out high performance computing. There are about
61 480 available GPUs, including models A10, A40, and
62 A100, and around 53000 CPU cores [1]. Each user is
63 allowed to store about 10TB of data across different
64 clusters.

65 2.1.2 Scheduling system

66 To compute a job, it must be submitted to a queue
67 using the *qsub* command. MetaCentrum's scheduler
68 then gives you access to resources based on their spe-
69 cial metric **fairshare** and the resources you requested.
70 Fairshare is a metric that evaluates how well you used
71 the resources available to you. Having low fairshare
72 means that you have a lower priority compared to other
73 researchers, which results in longer queue times [4].
74 Fairshare ranking example can be seen in Figure 1.
75 When training, I always requested 2 CPUs, 1 GPU and
76 24/32/48GB of RAM and VRAM based on the size of
77 the Whisper model in use - small/medium/largev2.

2.1.3 Logging

78 Since the original script lacked any form of logging
79 except plain text to standard and error output, logging
80 to weights and biases was integrated. That is where all
81 the graphs in this paper come from. The chosen model
82 and other parameters are passed either as arguments or
83 as environment variables.

85 2.2 Dataset

86 In order to replicate the results, the same data set that
87 the original paper used and the same speaker embed-
88 dings had to be obtained. That being the Libri2Mix
89 data set and its x-vector embeddings [2] .

90 2.2.1 Speech mixtures

91 I obtained the speech mixtures using LibriMix GitHub
92 repository available [here](#). When generating, I made
93 sure to generate only 16kHz audio and only Libri2Mix.

94 2.2.2 Metadata and embeddings

95 Metadata used for training has been generated using
96 a script provided [here](#) and embeddings were down-
97 loaded from [here](#). The created metadata are called
98 *train-100.json*, *train-100-noisy.json*, *test.json* and *test-*
99 *noisy.json*.

100 2.3 Whisper

101 Whisper is a state-of-the-art, open source, multilin-
102 gual speech transcription model developed for single-
103 speaker ASR instances. It can also produce speech
104 timestamps, perform inverse text normalization, etc.
105 Its input is an 80-mension log-Mel spectrogram and
106 its output is transcribed text [2].

107 2.4 Prompt tuning

108 Prompt tuning is a method of extending a model's capa-
109 bilities without changing the original weights (which
110 would be fine-tuning). It is a lightweight technique
111 that prepends a small number of trainable tokens to the
112 input called soft prompts. In our case, these tokens are
113 trained on target speaker embeddings [2].

114 2.4.1 Deep prompting

115 Deep prompting is an extension of prompt tuning, as
116 now prompts are inserted not only at the input but also
117 in multiple layers within the model. This results in
118 lower control over the model, but a possibility of better
119 results [2].

120 2.4.2 Soft prompt reparametrization

121 It has been reported that directly optimizing soft prompts
122 may make the training process unstable, which is why
123 a two-layer multilayer perceptron is implemented to
124 make it more stable [2].

3. Validation

To validate the claimed results, I had to train the model myself and compare the results. The code is public, so all I needed to do was to run it. The first training run unfortunately crashed due to a type error where the script was expecting a tuple in place of a list, so that had to be fixed.

3.1 Training

In training, the script first runs a hard coded ten epochs of training on the clean data and only one on noisy data. The learning rate is scheduled using a step scheduler that changes the learning rate from 10^{-4} to 10^{-5} after five epochs. The suggested prompt length is sixteen, and my experiments confirm that [3]. In particular, the proposed script functions the best with a batch size of one, which is quite unusual. The script also offers the ability to choose whether to use deep prompting, reparametrization, both, or neither. Training is performed on metadata *train-100.json* and *train-100-noisy.json* created in Section 2.

3.2 Evaluation

Evaluation is performed on metadata that come from the same dataset as the metadata used for training. The metadata is called *test.json* and *test-noisy.json* and was created in Section 2. The evaluation is performed at the last checkpoint created in the training, and its output is the word error rate.

Table 1. Performance under different settings claimed by Hao Ma and his team [2]. DP stands for deep prompting and MLP stands for multi layer perceptron. L/M/S stands for Whisper LargeV2/Medium/Small respectively

DP	MLP	test-clean			test-both		
		L	M	S	L	M	S
×	×	17.59	17.73	31.03	37.26	39.38	56.30
×	✓	15.92	14.61	24.27	32.25	32.34	45.72
✓	×	14.82	13.89	24.62	30.19	29.81	45.46
✓	✓	14.78	13.54	23.08	30.71	30.72	44.16

4. Experiments

In my attempt to further improve this approach, I chose to use only the medium model, as it proved to be the best one in Section 3. The second choice I made was to only experiment while using both deep prompting and reparametrization as it, in total, proved the best

Table 2. Performance under different settings achieved by me

DP	MLP	test-clean			test-both		
		L	M	S	L	M	S
×	×	17.43	16.06	30.10	37.82	36.57	51.54
×	✓	16.53	14.72	23.87	33.47	31.51	44.74
✓	×	16.12	14.64	23.86	31.66	30.75	44.47
✓	✓	16.80	14.62	23.07	32.32	31.04	43.57

in seven out of twelve cases in Tables 1 and 2. Then I and my supervisor came up with some ideas that I would like to try:

1. Training more epochs with a larger batch size
2. Using different schedulers for training
3. Mixing clean and noisy data into one dataset

4.1 More epochs and a larger batch size

In this attempt, I increased the total number of training epochs on clean data from ten to twenty and kept the total number of epochs on noisy data. Then I increased the batch size from one to sixteen. This has proven to be a bad approach as the word error rate on clean data increased to **17.63** and to **34.28** on noisy data which is way higher than WER obtained in Section 3.

4.2 Using different schedulers

In this approach, I set the batch size at one and the clean training epochs at ten. Then I increased the training epochs on noisy data to five. After that, I chose five different commonly used schedulers to determine the best one. The hypothesis that using different learning rate schedulers can affect the final performance of a speech model by better adapting the optimization process. The results can be seen in Table 3.

Scheduler	Clean	Both
Step	15.12	31.62
Linear	14.72	30.91
Cosine	15.57	29.96
Exponential	15.34	30.06
One cycle	15.89	32.48

Table 3. Performance of different schedulers on clean and noisy data

The results of individual schedulers on clean data turned out to be about the same for all of them and just a little worse than the results in Table 2. This could probably be attributed to longer training on noisy data.

185 However, on noisy data, we can see that the Cosine
 186 annealing scheduler managed to get below 30% WER
 187 which is about a 1% improvement on Table 2. Most
 188 importantly, it beats the step scheduler used in the orig-
 189 inal code by almost 2% on the same task, suggesting
 190 that it could be better for instances where a noisy en-
 191 vironment is expected. Also notable is that the linear
 192 scheduler beats the step one on both clean and noisy ut-
 193 terances, suggesting that it could be better for general
 194 use.

195 4.3 Mixing clean and noisy datasets

196 In this approach, I chose to mix the training data into
 197 one large file. Now instead of training on clean data
 198 and on noisy after that, the model trains on both at the
 199 same time. To make sure I get the most information
 200 from this approach, I also used five different sched-
 201 ulers. The hypothesis was that by having the model
 202 exposed to a broader range of acoustic conditions, the
 203 results would be more consistent. The results can be
 204 seen in Table 4.

Scheduler	Clean	Both
Step	19.49	33.73
Linear	18.93	33.74
Cosine	19.80	34.29
Exponential	18.20	32.65
One cycle	20.76	34.32

Table 4. Performance of different schedulers after training on mixed data

205 The results of this experiment are much worse than
 206 those in Table 2. The best result was achieved by the
 207 exponential scheduler in both clean and noisy environ-
 208 ments, but it was still about 3% worse than the previous
 209 results. It seems that this experiment has introduced
 210 some conflicting acoustic patterns that hindered the
 211 model’s learning capabilities. It is clear that training
 212 on a mixed set of data is not a good approach.

213 4.4 Results

214 In the end, I have managed to get results within close
 215 range of those claimed by Hao Ma, et al. in their pa-
 216 per [2]. To be exact, the average difference between
 217 mine and their results was 1.12%. When taking into
 218 consideration the training time required to get the re-
 219 sults, I can say that their solution is a strong one when
 220 it comes to environments without background noise,
 221 with WER around 14-15%, but in noisy environments,
 222 not so much as the WER is way higher at around 30-
 223 31%. The obtained results are shown in Table 2 and

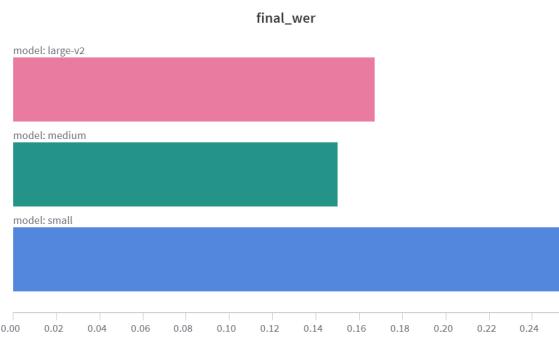


Figure 2. Graph showing performance by model on clean data

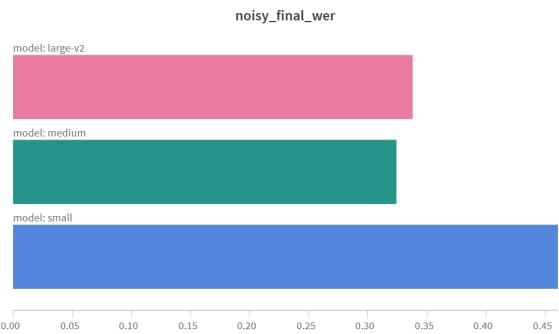


Figure 3. Graph showing performance by model on noisy data

those of the original article are shown in Table 1. What
 224 can be seen in the results is that using both yields
 225 the best results only about half of the time, the other
 226 half being deep prompting by itself. Further, what
 227 can be seen from the results is that the medium model
 228 achieves the best results in both clean and noisy envi-
 229 ronments. That can be best seen in Figures 2 and 3
 230 where the medium model beats the small one by a large
 231 margin and the large one by a little bit.
 232

233 5. Conclusions

This work has replicated the results claimed by Hao
 234 Ma, Zhiyuan Peng, Mingjie Shao, Jing Li, and Ju Liu.
 235 Prompt tuning is a fair approach when dealing with
 236 target speaker automatic speech recognition. This ap-
 237 proach avoids having to do expensive fine-tuning or
 238 building a model from scratch and achieves a much
 239 lower amount of resources and time needed than previ-
 240 ous state-of-the-art techniques [2].
 241

The best result on clean data is a word error rate
 242 of 14.62% and the best result on noisy data is a word
 243 error rate of 29.96% with a step scheduler as shown
 244 in Tables 2 and 3. Both have been obtained using
 245 Whisper medium and with both deep prompting and
 246 reparametrization enabled. Each of them used different
 247 schedulers, step and cosine, respectively, and different
 248

249 numbers of training epochs on noisy data. The most
250 consistent scheduler is the linear one, beating the step
251 one in the same number of training epochs in both
252 clean and noisy utterances, as well as beating the co-
253 sine one on clean data. However, all of this still falls
254 well short of the 6.9% WER on clean data and 15.9%
255 WER on noisy data reported by Alexander Polok's
256 team in the paper on Diarization-Conditioned Whisper
257 (DiCoW) [5].

258 Further work should be focused on implementing
259 support for multi-GPU training as that will cut training
260 time even more. The difference between using speaker
261 embeddings and diarization on noisy data suggests that
262 taking the path of diarization, such as DiCoW, is the
263 better choice for TS ASR.

264 Acknowledgements

265 I would like to thank my supervisor Alexander Polok
266 for leading me and helping me with this work.
267 This work was supported by the Ministry of Educa-
268 tion, Youth and Sports of the Czech Republic through
269 the e-INFRA CZ (ID:90140), provided within the pro-
270 gram Projects of Large Research, Development and
271 Innovations Infrastructures. Computational resources
272 were supplied by the project "e-Infrastructure CZ" (e-
273 INFRA CZ) and MetaCentrum.

274 References

- 275 [1] CESNET. *Grid Service MetaCentrum*.
276 2024. Accessed: 2025-06-10. Available
277 at: [https://www.cesnet.cz/gimg/
278 default/8/0/7/807-Grid%20service%
279 20MetaCentrum.pdf](https://www.cesnet.cz/gimg/default/8/0/7/807-Grid%20service%20MetaCentrum.pdf).
- 280 [2] MA, H., PENG, Z., SHAO, M., LI, J. and LIU,
281 J. *Extending Whisper with Prompt Tuning to*
282 *Target-Speaker ASR*. 2023. Accessed: 2025-06-
283 08. Available at: [https://arxiv.org/abs/
284 2312.08079](https://arxiv.org/abs/2312.08079).
- 285 [3] MA, H., PENG, Z., SHAO, M., LI, J. and LIU, J.
286 Extending Whisper with Prompt Tuning to Target-
287 Speaker ASR. In: *IEEE ICASSP*. 2024.
- 288 [4] METACENTRUM. *Welcome to MetaCentrum Doc-
289 umentation*. 2025. Accessed: 2025-06-08. Avail-
290 able at: [https://docs.metacentrum.cz/
291 en/docs/welcome](https://docs.metacentrum.cz/en/docs/welcome).
- 292 [5] POLOK, A., KLEMENT, D., KOCOUR, M., HAN,
293 J., LANDINI, F. et al. *DiCoW: Diarization-
294 Conditioned Whisper for Target Speaker Auto-
295 matic Speech Recognition*. 2024. Available at:
296 <https://arxiv.org/abs/2501.00114>.